

CSC 225 - Summer 2019


Sorting III

Bill Bird

Department of Computer Science
University of Victoria


June 5, 2019

Comparison Sorting (1)



- ▶ General sorting algorithms cannot make any assumptions about the input data.
- ▶ Comparison sorting algorithms, such as Merge Sort, Quicksort and Heap Sort, only interact with the input data by asking questions of the form ' $a \leq b$ '.
- ▶ We assume that for all input items x and y , either $x \leq y$ or $y \leq x$ (or both). This property is called 'total ordering'.
- ▶ The sorting algorithm is not concerned with how the \leq operator works.

Comparison Sorting (2)



- ▶ Comparison sorting algorithms must make $\Omega(n \log n)$ queries of the form $a \leq b$ in the worst case.
- ▶ The proof of this lower bound is now part of CSC 226.

Comparison of Comparison Sorting Algorithms (1)

n	Max. Comparisons			
	Selection Sort	Merge Sort	Quicksort	Heap Sort
3	3	3	3	3
4	6	5	6	7
5	10	8	10	12
6	15	11	15	17
7	21	14	21	22
8	28	17	28	29
9	36	21	36	37
10	45	25	45	44
11	55	29	55	51
12	66	33	66	59
13	78	37	78	66
14	91	41	91	73
15	105	45	105	80

This table gives the maximum number of comparisons performed by each algorithm over all arrays of size $n = 3, \dots, 15$.

Comparison of Comparison Sorting Algorithms (2)

n	Average Comparisons			
	Selection Sort	Merge Sort	Quicksort	Heap Sort
3	3.0	2.7	2.7	3.0
4	6.0	4.7	4.8	6.5
5	10.0	7.2	7.4	11.0
6	15.0	9.8	10.3	15.1
7	21.0	12.7	13.5	19.8
8	28.0	15.7	16.9	25.8
9	36.0	19.2	20.6	32.5
10	45.0	22.7	24.4	38.6
11	55.0	26.3	28.5	45.2
12	66.0	30.0	32.7	51.8
13	78.0	33.8	37.0	59.0
14	91.0	37.7	41.6	65.5
15	105.0	41.7	46.2	72.5

This table gives the average number of comparisons performed by each algorithm over all arrays of size $n = 3, \dots, 15$.

Sorting Under Assumptions (1)

314		323		314		159
159		843		323		265
265		383		843		314
358		314		358		323
979	→	265	→	159	→	358
323		358		265		383
843		159		979		843
383		979		383		979


When assumptions can be made about the input data, the $\Omega(n \log n)$ bound can sometimes be broken.

Sorting Under Assumptions (2)

314		323		314		159
159		843		323		265
265		383		843		314
358		314		358		323
979	→	265	→	159	→	358
323		358		265		383
843		159		979		843
383		979		383		979


For example, when all input values are 3-digit integers, the Radix Sort algorithm requires $O(n)$ time.

Bucket Sort (1)



Bucket Sort is a simple algorithm for sorting sequences with a small number of possible values. In the above example, there are only five possible elements.


Bucket Sort (2)



Clover	Star	Pot of Gold	Rainbow	Heart

A **bucket** is created for each type of element.


Bucket Sort (3)





Clover	Star	Pot of Gold	Rainbow	Heart

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (4)






Clover	Star	Pot of Gold	Rainbow	Heart
				

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (5)







Clover	Star	Pot of Gold	Rainbow	Heart
		 		

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (6)








Clover	Star	Pot of Gold	Rainbow	Heart
		 		

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (7)









Clover	Star	Pot of Gold	Rainbow	Heart
		 		 

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (8)










Clover	Star	Pot of Gold	Rainbow	Heart
		 		 

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (9)











Clover	Star	Pot of Gold	Rainbow	Heart
		 		  

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (10)



Clover	Star	Pot of Gold	Rainbow	Heart
		 		  

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (11)












Clover	Star	Pot of Gold	Rainbow	Heart

The first part of bucket sort loops over the input array, adding each element to its respective bucket.


Bucket Sort (12)












Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

After all elements have been added to buckets, the array is cleared and used to store the sorted output.


Bucket Sort (13)












Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

To generate the sorted array, the buckets are visited in ascending order and each bucket's contents are appended to the array.


Bucket Sort (14)












Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

To generate the sorted array, the buckets are visited in ascending order and each bucket's contents are appended to the array.

Bucket Sort (15)












Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

To generate the sorted array, the buckets are visited in ascending order and each bucket's contents are appended to the array.


Bucket Sort (16)












Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

To generate the sorted array, the buckets are visited in ascending order and each bucket's contents are appended to the array.


Bucket Sort (17)



Clover	Star	Pot of Gold	Rainbow	Heart
		 	 	  

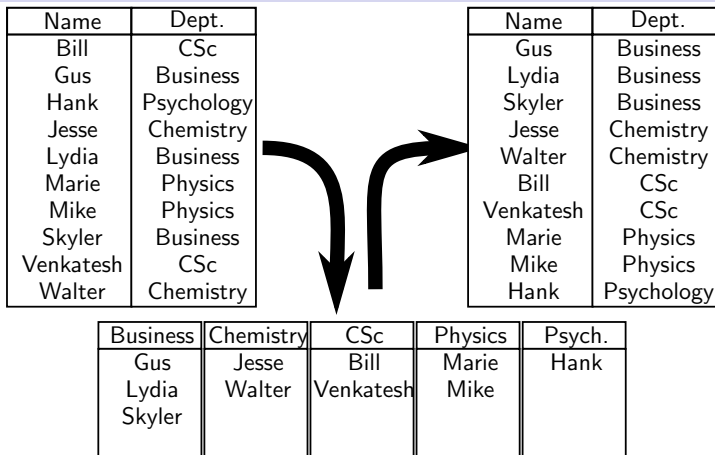
To generate the sorted array, the buckets are visited in ascending order and each bucket's contents are appended to the array.

Bucket Sort (18)



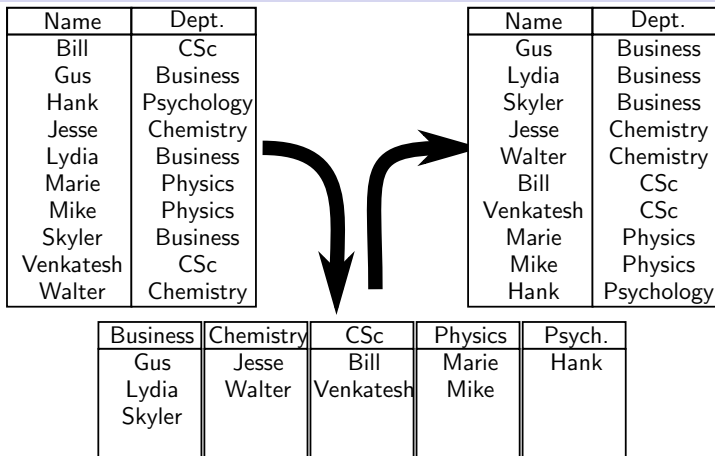
Bucket sort is very efficient when sorting on a key with a limited number of values (such as sorting employees by department in the example above).

Bucket Sort (19)



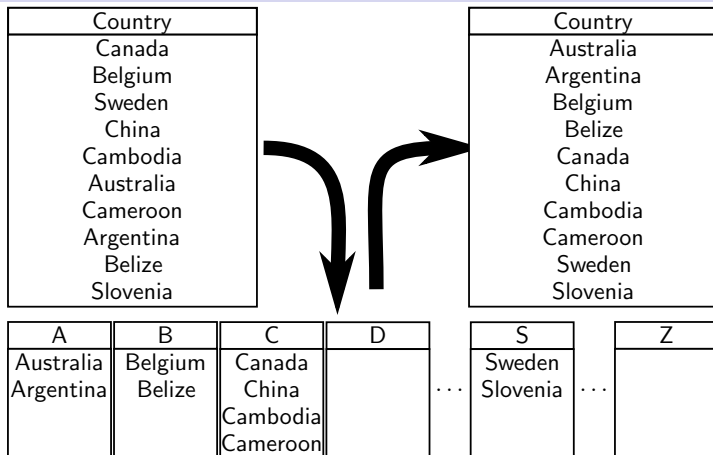
If there are b buckets, the running time of bucket sort on an array of size n is $O(n + b)$.

Bucket Sort (20)



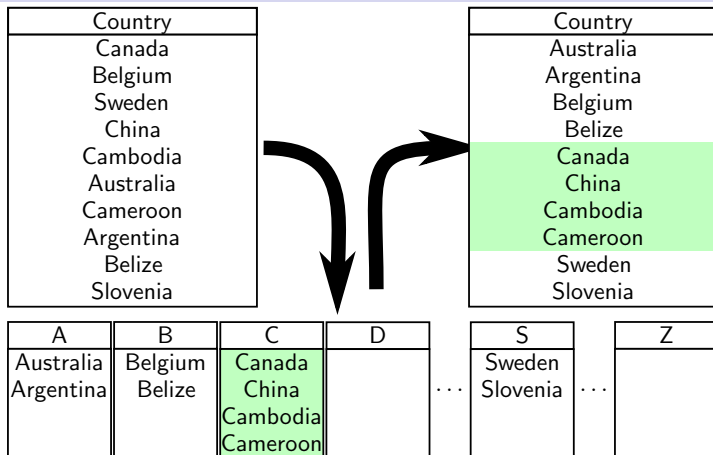
Bucket sort does not use pairwise comparisons at all, which is why the running time breaks the $\Omega(n \log n)$ bound.

Alphabetizing With Bucket Sort (1)




Using 26 buckets, a set of strings can be sorted by first letter using bucket sort.

Alphabetizing With Bucket Sort (2)




However, bucket sort does not alphabetize the strings.

Alphabetizing With Bucket Sort (3)



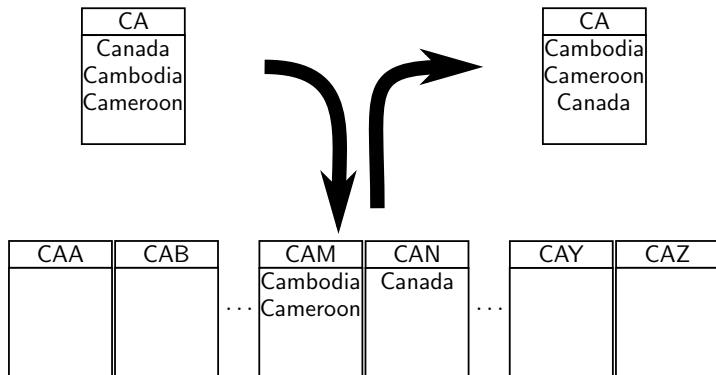
For example, the elements of the 'C' bucket are not in alphabetical order, since bucket sort puts them in the bucket in the order they were found.

Alphabetizing With Bucket Sort (4)




One solution is to recursively use bucket sort on each bucket and sort by the second letter...

Alphabetizing With Bucket Sort (5)




...then by the third letter, and continue until all buckets have size 1.

Alphabetizing With Bucket Sort (6)




If the strings in the list have maximum length k , sorting may require k passes.

Alphabetizing With Bucket Sort (7)




After the contents of each bucket are alphabetized, the resulting list is in alphabetical order.

Alphabetizing With Bucket Sort (8)



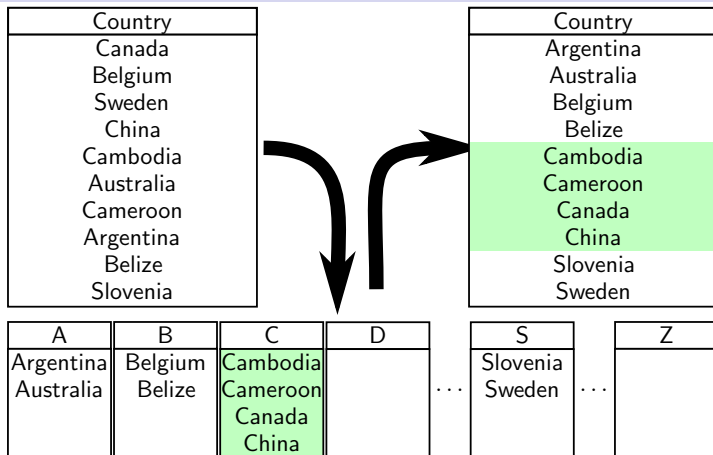
In a list of n strings with maximum length k , each string will be subjected to at most k sorting passes. Therefore, the list can be alphabetized in $O(nk)$ time.

Alphabetizing With Bucket Sort (9)



This variant of bucket sort is often called 'lexicographical sort' or 'dictionary sort'.

Alphabetizing With Bucket Sort (10)



When the maximum length k is bounded by a constant (e.g. 10 characters), lexicographical sort runs in $O(n)$ time.

Sorting Algorithms

Best Case	Running Time		Extra Space	Stable?
	Expected Case	Worst Case		

Selection Based

Heap Sort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(1)$	No
Insertion Sort	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	Yes
Selection Sort	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	No

Divide and Conquer

Merge Sort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n)$	Yes
Quicksort	$\Theta(n)$	$\Theta(n \log n)$	$\Theta(n^2)$	$\Theta(n)$	Yes

Other

Bubble Sort	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	Yes
Radix Sort ¹	$\Theta(dn + b)$	$\Theta(dn + b)$	$\Theta(dn + b)$	$\Theta(n + b)$	Yes

¹Integers only: d -digit values in base b

Sorting Algorithms

Best Case	Running Time		Extra Space	Stable?
	Expected Case	Worst Case		

Selection Based

Heap Sort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(1)$	No
Insertion Sort	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	Yes
Selection Sort	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	No

Divide and Conquer


Merge Sort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n)$	Yes
Quicksort	$\Theta(n)$	$\Theta(n \log n)$	$\Theta(n^2)$	$\Theta(n)$	Yes

Other

Bubble Sort	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	Yes
Radix Sort ¹	$\Theta(dn + b)$	$\Theta(dn + b)$	$\Theta(dn + b)$	$\Theta(n + b)$	Yes


¹Integers only: d -digit values in base b

Radix Sort (1)




Idea: Treat integers as strings of digits and use lexicographic sort.

Radix Sort (2)




By recursively sorting the buckets as before, this strategy does produce an effective sorting algorithm.

Radix Sort (3)



The resulting algorithm is **radix sort**. It can be performed in any base (base 10 is used here).

Radix Sort (4)



The running time of radix sort is $\Theta(d(n + b))$, where

d = number of digits

n = number of elements in input list

b = base (i.e. number of buckets)

MSD Radix Sort (1)

```
1: procedure MSDRADIXSORT( $A, n, k$ )
2:   if  $k < 0$  then
3:     return
4:   end if
5:    $B \leftarrow$  Array of 10 empty lists.
6:   for  $i = 0, 1, \dots, n - 1$  do
7:      $\text{digit} \leftarrow k^{\text{th}}$  digit of  $A[i]$ 
8:     Add  $A[i]$  to the end of list  $B[\text{digit}]$ 
9:   end for
10:   $S \leftarrow$  Empty list.
11:  for  $j = 0, 1, \dots, 9$  do
12:    if  $\text{len}(B[j]) > 1$  then
13:      MSDRADIXSORT( $B[j], \text{length}(B[j]), k - 1$ )
14:    end if
15:    Append  $B[j]$  to the end of  $S$ 
16:  end for
17:  Copy  $S$  into  $A$ 
18: end procedure
```

This variant is called Most Significant Digit (MSD) Radix Sort, since it starts with the most significant (leftmost) digit.

MSD Radix Sort (2)

```
1: procedure MSDRADIXSORT( $A, n, k$ )
2:   if  $k < 0$  then
3:     return
4:   end if
5:    $B \leftarrow$  Array of 10 empty lists.
6:   for  $i = 0, 1, \dots, n - 1$  do
7:      $\text{digit} \leftarrow k^{\text{th}}$  digit of  $A[i]$ 
8:     Add  $A[i]$  to the end of list  $B[\text{digit}]$ 
9:   end for
10:   $S \leftarrow$  Empty list.
11:  for  $j = 0, 1, \dots, 9$  do
12:    if  $\text{len}(B[j]) > 1$  then
13:      MSDRADIXSORT( $B[j], \text{length}(B[j]), k - 1$ )
14:    end if
15:    Append  $B[j]$  to the end of  $S$ 
16:  end for
17:  Copy  $S$  into  $A$ 
18: end procedure
```


To sort numbers with d digits, the initial call is MSDRADIX-SORT($A, \text{length}(A), d - 1$)

MSD Radix Sort (3)

```
1: procedure MSDRADIXSORT( $A, n, k$ )
2:   if  $k < 0$  then
3:     return
4:   end if
5:    $B \leftarrow$  Array of 10 empty lists.
6:   for  $i = 0, 1, \dots, n - 1$  do
7:      $\text{digit} \leftarrow k^{\text{th}}$  digit of  $A[i]$ 
8:     Add  $A[i]$  to the end of list  $B[\text{digit}]$ 
9:   end for
10:   $S \leftarrow$  Empty list.
11:  for  $j = 0, 1, \dots, 9$  do
12:    if  $\text{len}(B[j]) > 1$  then
13:      MSDRADIXSORT( $B[j], \text{length}(B[j]), k - 1$ )
14:    end if
15:    Append  $B[j]$  to the end of  $S$ 
16:  end for
17:  Copy  $S$  into  $A$ 
18: end procedure
```

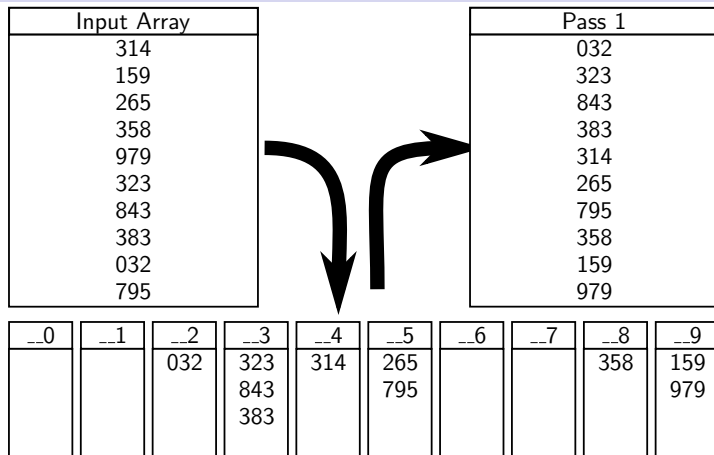
Note that digits are numbered starting at the right (so digit 0 is the least significant digit).

LSD Radix Sort (1)




An alternative is Least Significant Digit (LSD) Radix Sort. Instead of recursively calling bucket sort, LSD Radix Sort does a sequence of d bucket sorts over the whole array.

LSD Radix Sort (2)




In an array of d digit numbers, d passes are needed.

LSD Radix Sort (3)




In pass 1, the array is sorted by its least significant digit (digit 0). The output of pass 1 is used as the input of pass 2. After pass 1, the buckets can

LSD Radix Sort (4)




In pass 2, the output of pass 1 is sorted by digit 1.

LSD Radix Sort (5)



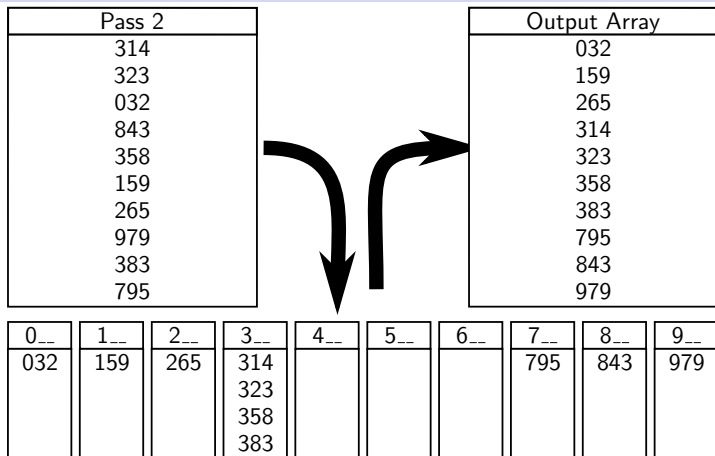
The bucket sort at each step must always insert values at the end of each bucket's list for LSD Radix Sort to work.

LSD Radix Sort (6)



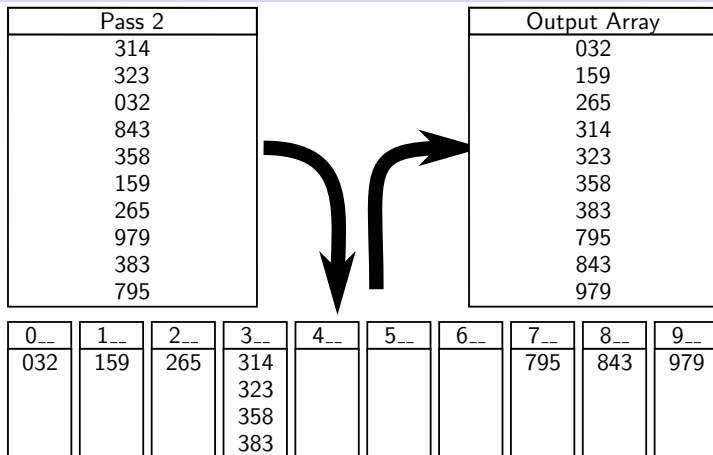
In pass 3, the output of pass 2 is sorted by digit 2.

LSD Radix Sort (7)



Since the values in the example above are 3 digits long, pass 3 is the final pass.

LSD Radix Sort (8)



Like MSD Radix Sort, the running time of LSD Radix Sort is $\Theta(d(n+b))$.

LSD Radix Sort (9)

```
1: procedure LSDRADIXSORT( $A, n, d$ )
2:   for  $k = 0, 1, \dots, d - 1$  do
3:      $B \leftarrow$  Array of 10 empty lists.
4:     for  $i = 0, 1, \dots, n - 1$  do
5:        $\text{digit} \leftarrow k^{\text{th}}$  digit of  $A[i]$ 
6:       Add  $A[i]$  to the end of list  $B[\text{digit}]$ 
7:     end for
8:      $S \leftarrow$  Empty list.
9:     for  $j = 0, 1, \dots, 9$  do
10:      Append  $B[j]$  to the end of  $S$ 
11:    end for
12:    Copy  $S$  into  $A$ 
13:  end for
14: end procedure
```

For sorting lists of integers, LSD Radix Sort tends to have less overhead in practice than MSD Radix Sort.

Radix Sort vs. Comparison Sorting Bound (1)

1111	1110	1100	1000	0000
1110	1100	1000	0000	0001
1101	1010	0100	1001	0010
1100	1000	0000	0001	0011
1011	0110	1101	1010	0100
1010	0100	1001	0010	0101
1001	0010	0101	1011	0110
1000	0000	0001	0011	0111
0111	→ 1111	→ 1110	→ 1100	→ 1000
0110	1101	1010	0100	1001
0101	1011	0110	1101	1010
0100	1001	0010	0101	1011
0011	0111	1111	1110	1100
0010	0101	1011	0110	1101
0001	0011	0111	1111	1110
0000	0001	0011	0111	1111

The number of digits in the base b representation of n is $\lceil \log_b n \rceil$. Therefore, sorting n integers in the range $1, \dots, n$ with radix sort requires $\Omega(n \log_b n)$ time.

Radix Sort vs. Comparison Sorting Bound (2)

1111	1110	1100	1000	0000
1110	1100	1000	0000	0001
1101	1010	0100	1001	0010
1100	1000	0000	0001	0011
1011	0110	1101	1010	0100
1010	0100	1001	0010	0101
1001	0010	0101	1011	0110
1000	0000	0001	0011	0111
0111	→ 1111	→ 1110	→ 1100	→ 1000
0110	1101	1010	0100	1001
0101	1011	0110	1101	1010
0100	1001	0010	0101	1011
0011	0111	1111	1110	1100
0010	0101	1011	0110	1101
0001	0011	0111	1111	1110
0000	0001	0011	0111	1111

Good Assignment Question: Give an algorithm, based on Radix Sort, which sorts an array of n integers in the range $[1, n^2]$ in $O(n)$ time.

Sorting Algorithm Summary (1)

None of the standard sorting algorithms covered in this course is 'always' the best choice. Depending on the application, and what assumptions can be made about the input data, any of the covered algorithms¹ could be the best choice.

The sorting algorithm currently used for general purpose comparison sorting by Python and Java is **Timsort**, which uses a combination of Merge Sort and Insertion Sort.

¹Except possibly Bubble Sort.

Sorting Algorithm Summary (2)

Selection Sort (worst case $\Theta(n^2)$):

- ▶ Tends to be surprisingly fast on small arrays.
- ▶ Predictable branching structure (which can speed up pipelined processors).
- ▶ Very easy to implement and likely to require few instructions.

Insertion Sort (worst case $\Theta(n^2)$):

- ▶ On nearly-sorted arrays, insertion sort is worst-case $\Theta(n)$
- ▶ Insertion sort is often the fastest algorithm on very small arrays.

Sorting Algorithm Summary (3)

Quicksort (worst case $\Theta(n^2)$):

- ▶ Expected case is $\Theta(n \log n)$.
- ▶ Tends to be faster in practice than Merge Sort.
- ▶ When pivots are chosen randomly, the likelihood of worst-case behavior is negligible.

Merge Sort (worst case $\Theta(n \log n)$):

- ▶ Guaranteed $\Theta(n \log n)$ running time on all inputs.
- ▶ Requires $\Theta(n)$ extra space.
- ▶ Can be used when the input sequence is too big to fit into memory (and is instead stored on disk), since all sequence processing is sequential.

Sorting Algorithm Summary (4)

Heap Sort (worst case $\Theta(n \log n)$):

- ▶ Guaranteed $\Theta(n \log n)$ running time on all inputs.
- ▶ Using `HEAPIFY`, Heap Sort can be implemented as an in-place algorithm.
- ▶ Very high overhead compared to Merge Sort and Quicksort.
- ▶ Heap Sort is not stable.

Radix Sort (worst case $\Theta(dn + b)$):

- ▶ $\Theta(n)$ worst-case performance on inputs with a constant number of digits.
- ▶ Relatively low overhead.
- ▶ No need for comparable data (Radix Sort performs zero comparisons).