

CSC 225 - Summer 2019

Dictionaries

Bill Bird

Department of Computer Science
University of Victoria

June 26, 2019

Dictionaries

A **Dictionary** data structure stores a collection of elements, each with a **key**. Dictionaries have three core operations:

1. **INSERT**(e): Add the element e to the collection.
2. **FIND**(k): Search for an element with key k and return it if it was found.
3. **REMOVE**(k): Remove key k from the dictionary.

In practice, keys are often a single component of a compound data type (for example, if a set of student records were stored in a dictionary, the student number or last name might be the key). In this course, for simplicity, keys are often the only data stored in dictionaries.

Dictionary Implementations

INSERT		FIND	
Expected Case	Worst Case	Expected Case	Worst Case

Sequence Based

Unsorted Array

Sorted Array

Unsorted List

$\Theta(1)$	$\Theta(n)$	$\Theta(n)$	$\Theta(n)$
$\Theta(n)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(\log n)$
$\Theta(1)$	$\Theta(1)$	$\Theta(n)$	$\Theta(n)$

Tree Based

'Normal' Binary Search Tree

2-3 Tree

Red/Black Tree

AVL Tree

B-Tree


$\Theta(\log n)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(n)$
$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$
$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$
$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$
$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$

Table Based

Hash Table


$\Theta(1)$	$\Theta(n)$	$\Theta(1)$	$\Theta(n)$
-------------	-------------	-------------	-------------

Binary Search Trees (1)




A *binary search tree* is a binary tree in which an in-order traversal produces a sorted sequence.

Binary Search Trees (2)




Alternatively, a binary search tree can be defined as a binary tree in which, for every node v , no nodes in the left subtree of v are greater than v and no nodes in the right subtree of v are less than v .

Binary Search Trees (3)




The subtrees of any node in a binary search tree are also binary search trees.

Binary Search Trees (4)



► Is this a binary search tree?

Binary Search Trees (5)




- ▶ Is this a binary search tree?
- ▶ **No.** The two highlighted nodes appear out of order in an in-order traversal.

BST Find (1)

```
1: procedure BSTFIND( $k$ , node)
2:   if node = null then
3:     return NOT_FOUND
4:   end if
5:   if node.value >  $k$  then
6:     return BSTFIND( $k$ , node.left)
7:   else if node.value <  $k$  then
8:     return BSTFIND( $k$ , node.right)
9:   else
10:    return node
11:  end if
12: end procedure
```


- ▶ The FIND Operation in a binary search tree can be implemented with a recursive search starting at the root.

BST Find (2)




- **Exercise:** Search for the key 11 in the binary search tree above.

BST Find (3)




- The initial call is `BSTFIND(11, root)`

BST Find (4)




- ▶ Since $11 > 10$, the search goes to the right.

BST Find (5)




- ▶ Since $11 < 17$, the search goes to the left.

BST Find (6)




- ▶ In the worst case, the BST_{FIND} algorithm traces a path from the root to a leaf.
- ▶ The total running time is $O(h)$ in a tree of height h .

BST Find (7)




- **Exercise:** Search for the key 18 in the binary search tree above.

BST Find (8)




- ▶ Since $18 > 10$, the search goes to the right.

BST Find (9)




- ▶ Since $18 > 17$, the search goes to the right.

BST Find (10)




- ▶ Since $18 < 19$, the search goes left.
- ▶ Since 19 has no left child, the search terminates unsuccessfully.

BST Insert (1)



- ▶ The INSERT procedure starts by searching for the correct location to insert the provided value.
- ▶ To insert the value 18, the search reaches the state above.

BST Insert (2)



- ▶ A new node is then created and inserted into the correct position.

BST Insert (3)

```
1: procedure BSTINSERT(node, new_value)
2:   if node = null then
3:     new_node  $\leftarrow$  New node containing new_value.
4:     return new_node
5:   end if
6:   if node.value  $\geq k$  then
7:     node.left  $\leftarrow$  BSTINSERT(node.left, new_value)
8:   else if node.value  $< k$  then
9:     node.right  $\leftarrow$  BSTINSERT(node.right, new_value)
10:  end if
11:  return node
12: end procedure
```

- Since the insertion itself is $\Theta(1)$, the BSTINSERT algorithm requires $\Theta(h)$ time.

BST Insert (4)

```
1: procedure BSTINSERT(node, new_value)
2:   if node = null then
3:     new_node  $\leftarrow$  New node containing new_value.
4:     return new_node
5:   end if
6:   if node.value  $\geq k$  then
7:     node.left  $\leftarrow$  BSTINSERT(node.left, new_value)
8:   else if node.value  $< k$  then
9:     node.right  $\leftarrow$  BSTINSERT(node.right, new_value)
10:  end if
11:  return node
12: end procedure
```

- **Question:** What is the maximum height of a binary search tree on n nodes?


BST Insert (5)

- ▶ **Exercise:** Insert the sequence 12, 6, 10, 11 into an initially empty binary search tree.

12


- ▶ **Exercise:** Insert the sequence 12, 6, 10, 11 into an initially empty binary search tree.

BST Insert (7)




- **Exercise:** Insert the sequence 12, 6, 10, 11 into an initially empty binary search tree.

BST Insert (8)




- **Exercise:** Insert the sequence 12, 6, 10, 11 into an initially empty binary search tree.

BST Insert (9)




- **Exercise:** Insert the sequence 12, 6, 10, 11 into an initially empty binary search tree.

BST Insert (10)




- ▶ The resulting tree contains 4 nodes and has height 3.
- ▶ In general, the height of a binary search tree is $\Theta(n)$ in the worst case.

BST Insert (11)




- ▶ As a result, the `FIND` and `INSERT` operations require $\Theta(n)$ time in the worst case.

BST Insert (12)




- ▶ The expected case running time of `FIND` and `INSERT` is $\Theta(\log n)$, since most sequences produce a relatively balanced tree.

BST Remove (1)




- **Exercise:** Delete the key 10 from the binary search tree above.

BST Remove (2)




- ▶ The deletion will leave a 'hole' in the tree which must be filled by a new element.

BST Remove (3)




- ▶ The element which replaces the root must preserve the binary search tree.
- ▶ The only elements which qualify are the next element (or previous element) in an in-order traversal.

BST Remove (4)




- After the replacement element is moved, a new hole is created.

BST Remove (5)




- Filling in the new hole is equivalent to recursively removing node 11.

BST Remove (6)




- ▶ The element 12 is moved to fill the hole created by 11 and the removal continues recursively.

BST Remove (7)



- ▶ Eventually, the removal process will end at a leaf, which can be deleted without negative consequences.

BST Remove (8)



- ▶ Eventually, the removal process will end at a leaf, which can be deleted without negative consequences.

Self-Balancing trees

To guarantee that both `FIND` and `INSERT` are $\Theta(\log n)$, it is necessary to ensure that the height of the tree is $\Theta(\log n)$. Some variants of the basic binary search tree, such as AVL trees and red/black trees, add a 'rebalancing' step to the `INSERT` method (after the normal binary search tree insertion) which maintains a $\Theta(\log n)$ height.

Self-balancing trees are covered in CSC 226.

2-3 Trees (1)

The 'rebalancing' approach used by AVL trees and red/black trees adjusts the tree after an insertion to compensate for any imbalances created by adding a node.

Odd Thought: What if, instead of rebalancing the tree after adding a node, we never add nodes at all? Instead, just cram new values into existing nodes, and split a node into multiple nodes if it gets too full.

Using nodes with maximum capacity 2 results in a data structure called a **2-3 tree**.


2-3 Trees (2)

- ▶ **Example:** Insert the sequence 31, 4, 15, 9, 2, 6, 5, 3 into an initially empty 2-3 tree.

31


- ▶ Create a new node to insert the first element. This is the only INSERT operation that directly creates a node.

2-3 Trees (4)




- ▶ To insert the second element, add it to the root node (which can contain up to 2 elements).

2-3 Trees (5)




- ▶ To insert 15, we first put it in the root node.
- ▶ However, the root is now too full, so it is split.

2-3 Trees (6)




- ▶ Splitting an over-full node produces two single nodes (for the smallest and largest elements in the over-full node) and elevates the middle element to the level above (which creates a new root).

2-3 Trees (7)




- ▶ To insert 9, use the `FIND` algorithm until the search terminates at a leaf.
- ▶ The element 9 is added to the node containing the 4.

2-3 Trees (8)




- ▶ Inserting 2 into the node with 4 and 9 makes the node too full, so it is split into three pieces: The 2 and 9 become single nodes, and the 4 is pushed upwards into the parent node.

2-3 Trees (9)




- ▶ A node with 2 keys a and b can have 3 children, corresponding to the elements less than a and b , elements between a and b , and elements greater than a and b .

2-3 Trees (10)




- The element 6 is inserted into the node with 9.

2-3 Trees (11)




- ▶ The element 5 is inserted into the node with 6 and 9.
- ▶ The resulting node is too full, so it is split, producing single nodes for 5 and 9, and pushing the 6 upwards.

2-3 Trees (12)




- Now the root is too full, so it is split, producing single nodes for 4 and 15 (which each take two of the child nodes), and pushing the 6 further upwards.

2-3 Trees (13)




- ▶ The resulting tree contains 6 at the root.
- ▶ Every level of a 2-3 tree is always full, and there are never any gaps.

2-3 Trees (14)




- The element 3 is inserted into the node containing 2.

2-3 Trees (15)




- ▶ The height of a 2-3 tree is always $\Theta(\log n)$, so FIND requires $\Theta(\log n)$ time.

2-3 Trees (16)



- ▶ The insertion process may require splitting one or more nodes, but the sequence of splits will progress through a single path from a leaf to the root, requiring $\Theta(\log n)$ time.

2-3 Trees (17)



- ▶ 2-3 trees are related to red/black trees, and are covered in greater detail in CSC 226.
- ▶ 2-3 trees are also a special case of B-trees, which are very important to database systems (covered in CSC 370).